

Member name:

**Sokdany Monyreach
Mey Chansopheaktra,
Thang Saoly,
Keav Sokhamuny,**

Rules for every challenge

- **Do all setup (filling, seeding indices, printing) outside the timed block.**
- **Put only the target work between t_0 and t_1 .**
- **Repeat ≥ 5 trials and average. Try multiple n (e.g., $1e3$, $1e4$, $1e5$) where it makes sense.**
- **Hand in: your code, a small table of times, and a 2–4 sentence explanation per part.**

Part A — Circular & Doubly Linked List mastery

A1 — CSLL: tail-to-head wrap vs manual reset

- **Implement traversal of n nodes in two ways:**
 - (i) CSLL with $\text{tail} \rightarrow \text{next} = \text{head}$ and a single loop of n steps;
 - (ii) Non-circular SLL that restarts at head whenever you hit `nullptr`.
- **Predict which is faster and why (branching, cache/predictability).**
- **A cache miss is typically much slower than a branch misprediction because the penalty of waiting for an external resource (RAM) is far greater than the penalty of a small pipeline rollback.**
- **Measure and explain.**

The key difference comes down to the source of the delay:

- **Cache Miss $\$l$ to $\$$ RAM:** Delay is caused by data unavailability from a slow, distant source (RAM). The penalty is in the hundreds of cycles.
- **Branch Misprediction:** Delay is caused by incorrect control flow requiring a pipeline reset, which is a fast, localized operation. The penalty is in the tens of cycles.

Therefore, although both are significant performance concerns, the cache miss to main memory is much slower because it forces the CPU to wait for a fundamental resource (data) from the slowest part of the memory hierarchy.

•

A2 — CSLL deletion with/without predecessor

- **Case 1:** Delete a given node when you also have its predecessor ($O(1)$).
- **Case 2:** Delete the same node when you only have the node pointer (must find predecessor).
- **Predict cost difference; measure for random positions; explain the curve.**

A3 — Rotate- k on CSLL vs SLL

- **Implement “rotate right by k ”:** in CSLL it’s pointer moves; in SLL it’s find-break-relink.
 1. In CSLL:
 - ❖ The list forms a circle because the last node points back to the head.
 - ❖ To rotate right by k , compute $k = k \% n$ for effective rotations.
 - ❖ Move the head pointer forward by $(n - k)$ steps.

- ❖ Since it's circular, no nodes need to be detached or relinked.
- ❖ Just update the head pointer to the new node at $(n - k)$ steps ahead.

2. In SLL:

- ❖ The last node points to nullptr, so the list is linear.
- ❖ Compute $k = k \% n$.
- ❖ Traverse to the $(n - k)$ th node (this will be the new tail).
- ❖ Find the last node (original tail).
- ❖ Break the link after the new tail by setting its **next** to **nullptr**.
- ❖ Link the original last node's **next** to the original head.
- ❖ Set the new head pointer to the node after the new tail.

- **Test for multiple k (small, $n/2$, $n-1$).**

- When k is small, both methods are fast.
- When $k = n/2$:
 - CSLL rotation still takes $O(k)$ pointer moves.
 - SLL rotation requires traversal and relinking, which is $O(n)$.
- When $k = n - 1$:
 - CSLL still just moves pointer.
 - SLL requires nearly full list traversal and relinking.

- **Decide** which wins and under what k ranges.

- CSLL wins for all k because rotation is just pointer moves require no relinking. But
- SLL rotation is requiring traversal and relinking.

A4 — DLL vs SLL: erase-given-node

- **Build DLL with prev.** Given a pointer to any node, erase it.
- **Compare to SLL** erase with known predecessor and SLL erase without predecessor.
- **Predict → measure → explain $O(1)$ vs $O(n)$ and constant-factor hits.**

A5 — Push/pop ends: head-only vs head+tail

- Implement `push_front/pop_front` and `push_back/pop_back` on:
 - (i) SLL with head only, (ii) SLL with head+tail, (iii) DLL with head+tail.
- **Benchmark each op for random mixes (e.g., 70% `push_back`, 30% `pop_front`).**
- **Explain why tail changes the story.**

A6 — Memory overhead audit

- **For the same logical dataset size n , allocate SLL, CSLL, and DLL nodes.**
- **Report bytes per node (pointer count), total bytes, and measured allocation time.**
- **Discuss the time-space trade-off you'd choose for frequent middle deletions.**

Part B — Real-world use cases

B1 — Recent Items Tray (add/remove at the same end)

- **Build a “recent items” tray** where the most recently added item is always the next one removed.
 - >This is created by putting what's add/remove on record, creating another and always putting it to the head pointer when inserting new or put into what just removed.
- **Implement with:**

- A) Singly linked nodes adding/removing at the front.
- B) Doubly linked nodes adding/removing at the front
- **Predict which is faster** (pointer count vs rewiring), measure adds/removes (mixed workload), explain whether the second pointer helps here.
->Total pointer ops (mixed add/remove workload):
 If you do 50% adds + 50% removes:

 Singly: ~4 pointer writes per full add/remove cycle

 Doubly: ~7 pointer writes per cycle
 So doubly linked list performs **~75% more pointer updates** for the same logical workload.

B2 — Editor Undo History

- **Implement an Undo history where each new action is “placed on top,” and Undo removes the last action added.**
- **Version A: Singly linked front-add/front-remove.**
- **Version B: Dynamic array (grow by doubling).**
- **Workload: 80% add actions, 20% undo actions.**

->

Metric	Singly Linked List	Dynamic Array
Add / Undo time	O(1) always	Amortized O(1), but realloc spikes
Memory locality	Poor (heap scattered)	Excellent (contiguous)
Memory overhead	1 pointer per node	Some unused capacity after growth
Realloc spikes	None	Yes (occasional copies)
Cache performance	Low	High
Predictability	Constant	Occasional stalls

- **Predict throughput and memory spikes (reallocs) vs constant-time links; measure and justify which design you’d ship.**
->Ship Version B (Dynamic Array)
 Because 80% adds, 20% undos → append/pop pattern fits a stack perfectly, and amortized O(1) realloc cost is negligible compared to improved cache performance.